

## Создание таблицы и вставка значений

```
CREATE TABLE products (
    id integer,
    created_at timestamp NOT NULL,
    active boolean NOT NULL,
    sort integer NOT NULL,
    price numeric(10,2) NOT NULL,
    code char(20) NOT NULL,
    name char(100) NOT NULL,
    description char(255) NOT NULL
);

INSERT INTO products (id, created_at, active, sort, price, code, name,
description)
SELECT
    gs AS id,
    timestamp '2016-01-01' + (random() * (365*8) * interval '1 day') AS
created_at,
    (random() < 0.70) AS active,
    (floor(random()*10000))::int AS sort,
    round((random()*10000)::numeric, 2) AS price,
    (array[
        'electronics', 'books', 'clothes', 'tools', 'sport',
        'beauty', 'games', 'home', 'kids', 'auto'
    ])[floor(random()*10)+1]::char(20) AS code,
    left(md5(random()::text), 60)::char(100) AS name,
    left(md5(random()::text) || md5(random()::text), 200)::char(255) AS
description
FROM generate_series(1, 10000) AS gs;
```

## Запрос №1

```
1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT * FROM products
3 WHERE id = 765432;
4
```

Data Output Messages Notifications

	id	created_at	active	sort	price	code	name	description
1	765432	2021-09-17 17:39:05.408671	false	271	8842.69	auto	3e53128c27b4c489184903b53d48e3...	7ebccdb19ed5d81bd11695dc75...

```
1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT * FROM products
3 WHERE id = 765432;
4
```

Data Output Messages Notifications

	QUERY PLAN
1	text
1	Gather (cost=1000.00..61764.47 rows=1 width=404) (actual time=65.162..77.545 rows=1 loops=1)
2	Workers Planned: 2
3	Workers Launched: 2
4	Buffers: shared hit=16118 read=39438
5	-> Parallel Seq Scan on products (cost=0.00..60764.38 rows=1 width=404) (actual time=51.778..54.461 rows=0 loops=3)
6	Filter: (id = 765432)
7	Rows Removed by Filter: 333333
8	Buffers: shared hit=16118 read=39438
9	Planning:
10	Buffers: shared hit=30 dirtied=3
11	Planning Time: 0.379 ms
12	Execution Time: 77.559 ms

## B-tree index (primary key)

```
ALTER TABLE products
ADD CONSTRAINT products_pkey PRIMARY KEY (id);
```

```
1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT * FROM products
3 WHERE id = 765432;
4
```

Data Output Messages Notifications



### QUERY PLAN text

```
1 Index Scan using products_pkey on products (cost=0.42..8.44 rows=1 width=404) (actual time=0.024..0.024 rows=1 loops=1)
2   Index Cond: (id = 765432)
3   Buffers: shared hit=1 read=3
4 Planning:
5   Buffers: shared hit=15 read=1
6 Planning Time: 0.791 ms
7 Execution Time: 0.033 ms
```

## Запрос №2

```
1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3   FROM products
4 WHERE active = true
5 ORDER BY sort
6 LIMIT 100
7
```

Data Output Messages Notifications



	<b>id</b> [PK] integer	<b>code</b> character (20)	<b>name</b> character (100)	<b>price</b> numeric (10,2)
1	35018	kids	e1df00e82e8db26880b4ca4488fbcd3...	8046.82
2	27970	home	c823f83daab9c86507f167f624b59ea...	9579.26
3	553339	clothes	8c6b4ebc67a2464a3cd675f83ec1d2b...	9011.02
4	13142	tools	202c91d85a63382601a4180cfe097bf...	3197.09
5	815662	sport	3115938da0312e04cda830d4f2c1b14...	931.85

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true
5 ORDER BY sort
6 LIMIT 100
7

```

Data Output Messages Notifications

	QUERY PLAN	
1	text	
1	Limit (cost=71808.38..71820.05 rows=100 width=136) (actual time=127.013..131.447 rows=100 loops=1)	
2	Buffers: shared hit=16232 read=39436	
3	-> Gather Merge (cost=71808.38..139492.61 rows=580110 width=136) (actual time=127.012..131.442 rows=100 loops=1)	
4	Workers Planned: 2	
5	Workers Launched: 2	
6	Buffers: shared hit=16232 read=39436	
7	-> Sort (cost=70808.36..71533.50 rows=290055 width=136) (actual time=107.505..107.508 rows=85 loops=3)	
8	Sort Key: sort	
9	Sort Method: top-N heapsort Memory: 69kB	
10	Buffers: shared hit=16232 read=39436	
11	Worker 0: Sort Method: top-N heapsort Memory: 68kB	
12	Worker 1: Sort Method: top-N heapsort Memory: 68kB	
13	-> Parallel Seq Scan on products (cost=0.00..59722.67 rows=290055 width=136) (actual time=0.124..81.807 rows=233396 loops=3)	
14	Filter: active	
15	Rows Removed by Filter: 99937	
16	Buffers: shared hit=16125 read=39431	
17	Planning Time: 0.075 ms	
18	Execution Time: 131.469 ms	

## Частичный индекс

CREATE INDEX idx\_products\_active\_sort ON products (sort) WHERE active = true;

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true
5 ORDER BY sort
6 LIMIT 100
7

```

Data Output Messages Notifications

	QUERY PLAN	
1	text	
1	Limit (cost=0.42..34.19 rows=100 width=136) (actual time=0.017..0.077 rows=100 loops=1)	
2	Buffers: shared hit=103	
3	-> Index Scan using idx_products_active_sort on products (cost=0.42..235041.13 rows=696133 width=136) (actual time=0.016..0.071 rows=100 loops=1)	
4	Buffers: shared hit=103	
5	Planning Time: 0.066 ms	
6	Execution Time: 0.088 ms	

### Запрос №3

```
1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true
5 ORDER BY price DESC
6 LIMIT 100
7
```

Data Output Messages Notifications

	<b>id</b> [PK] integer	<b>code</b> character (20)	<b>name</b> character (100)	<b>price</b> numeric (10,2)
1	18445	beauty	e2eaf985b4c7d8e96300484e16f166...	10000.00
2	424868	tools	91c459b639634b94045ddfc13546a...	9999.97
3	384558	books	568f63831292eefc21fb3e922de7b56...	9999.95
4	528362	clothes	36f95f31929a42e628921e5ce1c278...	9999.95
5	148632	auto	d94e2ae3c87b144260164e51290d81...	9999.95

```
1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true
5 ORDER BY price DESC
6 LIMIT 100
7
```

Data Output Messages Notifications

	QUERY PLAN
1	text
1	Limit (cost=71808.38..71820.05 rows=100 width=132) (actual time=111.748..116.244 rows=100 loops=1)
2	Buffers: shared hit=16258 read=39410
3	-> Gather Merge (cost=71808.38..139492.61 rows=580110 width=132) (actual time=111.747..116.239 rows=100 loops=1)
4	Workers Planned: 2
5	Workers Launched: 2
6	Buffers: shared hit=16258 read=39410
7	-> Sort (cost=70808.36..71533.50 rows=290055 width=132) (actual time=92.256..92.262 rows=84 loops=3)
8	Sort Key: price DESC
9	Sort Method: top-N heapsort Memory: 71kB
10	Buffers: shared hit=16258 read=39410
11	Worker 0: Sort Method: top-N heapsort Memory: 68kB
12	Worker 1: Sort Method: top-N heapsort Memory: 68kB
13	-> Parallel Seq Scan on products (cost=0.00..59722.67 rows=290055 width=132) (actual time=0.194..62.533 rows=233396 loops=3)
14	Filter: active
15	Rows Removed by Filter: 99937
16	Buffers: shared hit=16146 read=39410
17	Planning Time: 0.062 ms
18	Execution Time: 116.276 ms

## Частичный индекс

```
CREATE INDEX idx_products_active_price ON products (price DESC) WHERE active = true;
```

```
1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true
5 ORDER BY price DESC
6 LIMIT 100
7
```

Data Output Messages Notifications

	QUERY PLAN	
1	Limit (cost=0.42..34.95 rows=100 width=132) (actual time=0.038..0.443 rows=100 loops=1)	
2	Buffers: shared hit=23 read=80	
3	-> Index Scan using idx_products_active_price on products (cost=0.42..240316.98 rows=696133 width=132) (actual time=0.038..0.435 rows=100 loops=1)	
4	Buffers: shared hit=23 read=80	
5	Planning:	
6	Buffers: shared hit=15 read=1	
7	Planning Time: 0.718 ms	
8	Execution Time: 0.453 ms	

## Запрос №4

```
1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true AND id IN (1,500,70000,800000,10230,345234)
5 ORDER BY sort
6 LIMIT 100
7
```

Data Output Messages Notifications

	id [PK] integer	code character (20)	name character (100)	price numeric (10,2)
1	10230	home	4552d1f9368fccfd52118ae71fe17c6...	4118.68
2	800000	clothes	f3f5dccd088168228a4c3b63183cd3...	1463.94
3	1	tools	af77991928106d769a41becc3056cd...	25.52
4	70000	home	06d5eafc7ae062d117adbd3dcc96...	3017.89
5	345234	sport	05ef122b802d1a1ea045a0cd80bafb...	5897.75

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true AND id IN (1,500,70000,800000,10230,345234)
5 ORDER BY sort
6 LIMIT 100
7

```

Data Output Messages Notifications

	QUERY PLAN	text	SQL
1	Limit (cost=30.70..30.71 rows=4 width=136) (actual time=0.048..0.049 rows=5 loops=1)		
2	Buffers: shared hit=24		
3	-> Sort (cost=30.70..30.71 rows=4 width=136) (actual time=0.048..0.048 rows=5 loops=1)		
4	Sort Key: sort		
5	Sort Method: quicksort Memory: 26kB		
6	Buffers: shared hit=24		
7	-> Index Scan using products_pkey on products (cost=0.42..30.66 rows=4 width=136) (actual time=0.015..0.038 rows=5 loops=1)		
8	Index Cond: (id = ANY ('{1,500,70000,800000,10230,345234}'::integer[]))		
9	Filter: active		
10	Rows Removed by Filter: 1		
11	Buffers: shared hit=24		
12	Planning Time: 0.083 ms		
13	Execution Time: 0.061 ms		

Дополнительные индексы не требуются

## Запрос №5

```

1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true AND name LIKE '1f%'
5 LIMIT 100
6

```

Data Output Messages Notifications

	id [PK] integer	code character (20)	name character (100)	price numeric (10,2)
1	45179	auto	1fa8a0afed682973e42323a0900097...	4299.49
2	45663	clothes	1fc2ac330f0b8571e973008a37c881...	9784.11
3	46913	home	1fdfe5341d4af5c4d8b43af1f41a1f5...	2840.55
4	47031	books	1f774f4ed35becf0eb3c24420f92ad...	2603.88

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true AND name LIKE '1f%'
5 LIMIT 100
6

```

Data Output Messages Notifications

	QUERY PLAN	
	text	
1	Limit (cost=1000.00..61771.33 rows=70 width=132) (actual time=0.235..33.008 rows=100 loops=1)	
2	Buffers: shared hit=24 read=1994	
3	-> Gather (cost=1000.00..61771.33 rows=70 width=132) (actual time=0.234..32.998 rows=100 loops=1)	
4	Workers Planned: 2	
5	Workers Launched: 2	
6	Buffers: shared hit=24 read=1994	
7	-> Parallel Seq Scan on products (cost=0.00..60764.33 rows=29 width=132) (actual time=0.307..2.479 rows=34 loops=3)	
8	Filter: (active AND (name ~~ '1f%':text))	
9	Rows Removed by Filter: 12032	
10	Buffers: shared hit=24 read=1994	
11	Planning Time: 0.090 ms	
12	Execution Time: 33.059 ms	

## GIN

Подключение расширения триграмм (разбиение на 3 всевозможных слога):

CREATE EXTENSION IF NOT EXISTS pg\_trgm;

Создание GIN индекса по имени товара:

CREATE INDEX idx\_products\_active\_name\_trgm ON products USING gin ((name::text) gin\_trgm\_ops)  
WHERE active = true;

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE active = true AND name::text LIKE '1f%'
5 LIMIT 100
6

```

Data Output Messages Notifications

	QUERY PLAN	
	text	
1	Limit (cost=60.41..378.12 rows=100 width=132) (actual time=0.839..0.922 rows=100 loops=1)	
2	Buffers: shared hit=123	
3	-> Bitmap Heap Scan on products (cost=60.41..11119.96 rows=3481 width=132) (actual time=0.839..0.917 rows=100 loops=1)	
4	Recheck Cond: (((name)::text ~~ '1f%':text) AND active)	
5	Heap Blocks: exact=97	
6	Buffers: shared hit=123	
7	-> Bitmap Index Scan on idx_products_active_name_trgm (cost=0.00..59.54 rows=3481 width=0) (actual time=0.620..0.621 rows=2716 loops=1)	
8	Index Cond: ((name)::text ~~ '1f%':text)	
9	Buffers: shared hit=26	
10	Planning:	
11	Buffers: shared hit=1	
12	Planning Time: 0.095 ms	
13	Execution Time: 0.974 ms	

## Запрос №6

```
1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT *
3 FROM products
4 WHERE active = true AND created_at > '2023-12-01 00:00:00'
5 LIMIT 100
6
```

Data Output Messages Notifications

	<b>id</b> [PK] integer	<b>created_at</b> timestamp without time zone	<b>active</b> boolean	<b>sort</b> integer	<b>price</b> numeric (10,2)	<b>code</b> character (20)	<b>name</b> character (100)	<b>description</b> character (255)
1	326036	2023-12-16 01:13:15.84696	true	78	5.96	electronics	e80edef857e1c17ead45216f05fbe0...	0f8bdb03c93d0aafee8a60d2f7b57...
2	326060	2023-12-09 11:09:38.130211	true	3511	7901.76	electronics	ebe7a15964abc7733b22ffc912bcfa...	043da48349300352ded7e4c2b540...
3	326081	2023-12-06 22:50:08.363776	true	2151	4567.32	books	a0d1920ebbd3212b114730b13008...	ffb41c9ed3049212387113bdfb848...
4	326156	2023-12-21 20:25:20.006179	true	2642	9979.83	books	bc502e72866d289e7834283e10347...	8f7062664ec68b6cc2ad1980e80e...

```
1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT *
3 FROM products
4 WHERE active = true AND created_at > '2023-12-01 00:00:00'
5 LIMIT 100
6
```

Data Output Messages Notifications

	QUERY PLAN	text	lock
1	Limit (cost=0.00..951.97 rows=100 width=404) (actual time=0.012..2.754 rows=100 loops=1)		
2	Buffers: shared hit=15 read=784		
3	-> Seq Scan on products (cost=0.00..68056.00 rows=7149 width=404) (actual time=0.011..2.749 rows=100 loops=1)		
4	Filter: (active AND (created_at > '2023-12-01 00:00:00'::timestamp without time zone))		
5	Rows Removed by Filter: 14179		
6	Buffers: shared hit=15 read=784		
7	Planning:		
8	Buffers: shared hit=1		
9	Planning Time: 0.078 ms		
10	Execution Time: 2.792 ms		

## BRIN

```
CREATE INDEX idx_products_created_at_brin ON products USING brin (created_at);
```

```
1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT *
3 FROM products
4 WHERE active = true AND created_at > '2022-12-01 00:00:00'
5 LIMIT 100
6
```

Data Output Messages Notifications



QUERY PLAN  
text



```
1 Limit (cost=0.00..73.87 rows=100 width=404) (actual time=0.011..0.270 rows=100 loops=1)
2   Buffers: shared hit=17 read=46
3     -> Seq Scan on products (cost=0.00..68056.00 rows=92133 width=404) (actual time=0.009..0.264 rows=100 loop...
4       Filter: (active AND (created_at > '2022-12-01 00:00:00'::timestamp without time zone))
5       Rows Removed by Filter: 942
6     Buffers: shared hit=17 read=46
7 Planning:
8   Buffers: shared hit=2
9 Planning Time: 0.096 ms
10 Execution Time: 0.312 ms
```

## Запрос №7 (hash)

```
1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE code = 'auto' AND active = true
5
```

Data Output Messages Notifications



	<b>id</b> [PK] integer	<b>code</b> character (20)	<b>name</b> character (100)	<b>price</b> numeric (10,2)
1	516099	auto	f6d2bae2151a3cb402dd5e09ca08dd...	629.15
2	516100	auto	8dd1695e33fa74970a64e613f11822...	1450.53
3	516101	auto	a89a644f1bb8ad05702383079774b3...	2200.49

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE code = 'auto' AND active = true
5

```

Data Output Messages Notifications

	QUERY PLAN	
1	Seq Scan on products (cost=0.00..68056.00 rows=68523 width=132) (actual time=0.178..124.975 rows=70159 loops=...	
2	Filter: (active AND (code = 'auto'::bpchar))	
3	Rows Removed by Filter: 929841	
4	Buffers: shared hit=16346 read=39210	
5	Planning:	
6	Buffers: shared hit=1	
7	Planning Time: 0.087 ms	
8	Execution Time: 126.390 ms	

## HASH

CREATE INDEX idx\_products\_category\_hash ON products USING hash (code) WHERE active = true;

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, code, name, price
3 FROM products
4 WHERE code = 'auto' AND active = true
5

```

Data Output Messages Notifications

	QUERY PLAN	
1	Bitmap Heap Scan on products (cost=2607.05..61939.40 rows=68523 width=132) (actual time=6.866..142.639 rows=70159 loops=1)	
2	Recheck Cond: ((code = 'auto'::bpchar) AND active)	
3	Heap Blocks: exact=40592	
4	Buffers: shared read=40765	
5	-> Bitmap Index Scan on idx_products_active_code_hash (cost=0.00..2589.92 rows=68523 width=0) (actual time=3.629..3.629 rows=70159 loop...	
6	Index Cond: (code = 'auto'::bpchar)	
7	Buffers: shared read=173	
8	Planning:	
9	Buffers: shared read=1	
10	Planning Time: 0.109 ms	
11	Execution Time: 144.433 ms	

## Запрос №8 (jsonb)

CREATE TABLE api\_responses (
 id BIGSERIAL PRIMARY KEY,
 payload JSONB NOT NULL
);

```

INSERT INTO api_responses(payload)
SELECT
    case when random() < 0.1
        then jsonb_build_object('error','timeout','code',504,'req_id',gs)
        else jsonb_build_object('status','ok','req_id',gs)
    end
FROM generate_series(1, 10000) gs;

```

```

1 -- EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, payload
3 FROM api_responses
4 WHERE payload ? 'error'

```

Data Output Messages Notifications

	<b>id</b> [PK] bigint	<b>payload</b> jsonb
1	3012	{"code": 504, "error": "timeout", "req_id": 3012}
2	3039	{"code": 504, "error": "timeout", "req_id": 3039}
3	3050	{"code": 504, "error": "timeout", "req_id": 3050}
4	3055	{"code": 504, "error": "timeout", "req_id": 3055}
5	3060	{"code": 504, "error": "timeout", "req_id": 3060}

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, payload
3 FROM api_responses
4 WHERE payload ? 'error'

```

Data Output Messages Notifications

	QUERY PLAN text	Showing rows: 1 to 8
1	Seq Scan on api_responses (cost=10000000000.00..10000046275.00 rows=202020 width=55) (actual time=0.023..118.163 rows=200643 loops=...	
2	Filter: (payload ? 'error'::text)	
3	Rows Removed by Filter: 1799357	
4	Buffers: shared hit=16290 read=4985	
5	Planning:	
6	Buffers: shared hit=3 read=2 dirtied=1	
7	Planning Time: 0.412 ms	
8	Execution Time: 121.986 ms	

```

CREATE INDEX idx_api_responses_payload_gin ON api_responses USING GIN
(payload);

```

```

1 EXPLAIN (ANALYZE, BUFFERS)
2 SELECT id, payload
3 FROM api_responses
4 WHERE payload ? 'error'

```

Data Output Messages Notifications

	SQL	Showing rows: 1 to 11		
<b>QUERY PLAN</b>				
text				
1	Bitmap Heap Scan on api_responses (cost=1387.21..25187.46 rows=20200 width=55) (actual time=10.184..93.470 rows=200643 loops=1)			
2	Recheck Cond: (payload ? 'error'::text)			
3	Heap Blocks: exact=21275			
4	Buffers: shared hit=1 read=21308			
5	-> Bitmap Index Scan on idx_api_responses_payload_gin (cost=0.00..1336.70 rows=20200 width=0) (actual time=8.569..8.569 rows=200643 loop...			
6	Index Cond: (payload ? 'error'::text)			
7	Buffers: shared hit=1 read=33			
8	Planning:			
9	Buffers: shared hit=1			
10	Planning Time: 0.060 ms			
11	Execution Time: 97.575 ms			

## Выход индексов

```

1 SELECT
2   t.relname AS table_name,
3   i.relname AS index_name,
4   am.amname AS index_type, -- btree, gin, gist, hash и т.д.
5   pg_get_indexdef(i.oid) AS indexdef
6   FROM pg_class t
7   JOIN pg_index ix ON ix.indrelid = t.oid
8   JOIN pg_class i ON i.oid = ix.indexrelid
9   JOIN pg_am am ON am.oid = i.relam
10  JOIN pg_namespace n ON n.oid = t.relnamespace
11  WHERE n.nspname = 'public'
12  AND t.relname IN ('products', 'api_responses')
13  ORDER BY t.relname, i.relname;

```

Data Output Messages Notifications

	table_name name	index_name name	index_type name	indexdef text	
1	api_responses	api_responses_pkey	btree	CREATE UNIQUE INDEX api_responses_pkey ON public.api_responses USING btree (id)	
2	api_responses	idx_api_responses_payload_gin	gin	CREATE INDEX idx_api_responses_payload_gin ON public.api_responses USING gin (payload)	
3	products	idx_products_active_code_hash	hash	CREATE INDEX idx_products_active_code_hash ON public.products USING hash (code) WHERE (active = true)	
4	products	idx_products_active_name_trgm	gin	CREATE INDEX idx_products_active_name_trgm ON public.products USING gin (((name)::text) gin_trgm_ops) WHERE (active = true)	
5	products	idx_products_active_price	btree	CREATE INDEX idx_products_active_price ON public.products USING btree (price DESC) WHERE (active = true)	
6	products	idx_products_active_sort	btree	CREATE INDEX idx_products_active_sort ON public.products USING btree (sort) WHERE (active = true)	
7	products	idx_products_created_at_btree	brin	CREATE INDEX idx_products_created_at_btree ON public.products USING brin (created_at)	
8	products	products_pkey	btree	CREATE UNIQUE INDEX products_pkey ON public.products USING btree (id)	

